

Layered Controls and Service Disruption in a Synthetic Hospital Network: A Paired Monte Carlo Analysis with Observational Benchmarks

Ashish Agrawal¹ Mukil Dharanidharan¹ Naman Upadhyay¹

¹Independent Student Researchers, Naperville, Illinois, USA

Abstract Public ransomware studies document hospital disruption but cannot replay the same event under different controls. We built a Monte Carlo model of ransomware-like movement through a synthetic hospital network and compared a flat reference with basic segmentation, faster detection and isolation, and protected backups. Paired runs shared topology, entry point, and event-level random values. A time-step diagnostic changed the size of the initial effect, so we froze a five-minute protocol and ran 500 fresh pairs. Mean weighted service-hours lost fell from 193.4 to 33.9; the mean paired change was -159.5 hours (95% bootstrap interval, -166.1 to -152.7). Loss was lower in 451 pairs, tied in 48, and higher in one. All 128 joint stress settings favored the layered portfolio on their mean, but the median reduction was 28.7% and the fifth percentile was 3.1%.

We then added a non-calibrating bridge to real data. The public THREAT deposit yielded 149 attacked hospital-event records across 74 ransomware events; 30.9% carried an emergency-diversion flag, 49.0% carried a cancellation/delay flag, and 19 events involved multiple hospitals. A frozen CISA catalog contained 1,687 vulnerabilities exploited in the wild, including 352 tagged for known ransomware use. These observations did not identify a spread hazard or control effect. Instead, they exposed construct and structural failures: operational flags are not equivalent to binary modeled service-hours, the simulator cannot represent multi-hospital events, and its reference cases recovered faster than the roughly three-week hospital-volume pattern. The simulation result is therefore a directional within-model comparison, not a real-world effectiveness estimate.

Keywords healthcare cybersecurity; ransomware; Monte Carlo simulation; common random numbers; service continuity; model verification; observational benchmarking; external validity

1 Introduction

Counting infected computers misses the part of a hospital cyberattack that matters most to patients and staff: which services still work, and for how long. In 374 publicly reported attacks on U.S. healthcare delivery organizations from 2016 through 2021, 44.4% mentioned disruption to care and 41.7% mentioned electronic-system downtime. For incidents with a reported duration, the mean was 15.8 days [1]. Those numbers describe a selected set of visible attacks. They are not a hospital's probability of being attacked.

The operational record is becoming clearer. WannaCry was associated with lost appointments and reduced activity at affected National Health Service trusts [2]. During a month-long U.S. attack, two neighboring emergency departments handled more patients and ambulances and experienced longer waits and stroke-care delays across 19,857 visits [3]. Later work reported changes in cardiac-arrest incidence and outcomes at nearby, unaffected hospitals [4]. Using a national attack database linked to Medicare claims, Neprash et al. [5] estimated a 17% to 24% decline in hospital volume during the first attack week, with recovery over roughly three weeks. The studies differ in design and outcome, but they agree on one practical point: digital outages can become clinical and regional operating problems.

CISA, HHS, and NIST recommend preparation, segmentation, detection and response, identity protection, and recoverability [6–8]. Their guidance does not tell a modeler what number to use for a spread probability or

isolation-success rate, nor should it. We treated those quantities as assumptions to stress, not as measurements borrowed from guidance.

Our question was: *when the same synthetic incident is replayed under a flat reference posture and a feasible layered portfolio, does the layered portfolio preserve more modeled service availability across new scenarios, joint parameter changes, and different time steps?* We expected the answer to be yes. The useful test was whether that answer survived less favorable assumptions and whether its size remained stable. It did survive. Its size did not. The model is not an estimate of hospital risk, patient harm, or the causal effect of deployed controls.

1.1 Related empirical and simulation research

Downtime research helps connect unavailable software to work on the hospital floor. Among 80,381 patient-safety reports, Larsen et al. [9] found 76 that described EHR downtime. Laboratory and medication processes appeared often, and downtime procedures were not used consistently. A mixed-methods study of a roughly 48-hour outage found delays in staff accounts that were not fully visible in electronic timestamps [10]. Related AHRQ work used records, observation, interviews, and simulation to study laboratory and emergency-department contingencies [11]. This evidence is why our outcome is service availability rather than only the number of compromised nodes. It also makes clear what our binary service measure leaves out.

The closest simulation studies operate at different lay-

ers. A patient-flow model compared ransom payment and backup restoration across an emergency department, intensive care unit, and laboratory [12]. A hybrid process model examined how device outages and emergency plans affect trauma care [13]. Other work has modeled financial loss with bond percolation [14] and infrastructure interdependencies with agents [15]. Most recently, Arsalani et al. [16] combined data from two Alabama incidents with a regional EMS simulation and tested the role of hospital IT segmentation. That paper is especially relevant: it means we cannot claim to be the first to simulate ransomware, hospital resilience, or segmentation.

Our model stays inside one synthetic hospital. It connects digital propagation to seven service dependencies, then compares the same incident under two control portfolios. We do not model patient flow, financial loss, or regional demand. Our narrower contribution is an auditable pairing method, a joint test of 13 assumptions, and a documented time-step failure.

We kept a firm boundary between observations and assumptions. Incident studies justified multi-day checks and service-based outcomes [1, 2]. Work on neighboring hospitals showed that our single-hospital model misses regional spillover [3, 17]. Downtime research showed why manual workarounds and process delays matter [9–11]. Government guidance helped us choose relevant controls and failure modes [6–8]. None of those sources supplies a universal ransomware spread probability or a measured isolation-success rate, so we did not treat them as if they did.

2 Methods

2.1 Model scope

The Python model is defensive: it contains no malware, exploit code, scanning, or connection to a real network. A generated directed graph represents permitted communication or access paths among workstations, EHR, laboratory, pharmacy, imaging, medical devices, administration, identity, backup, and internet-facing/vendor functions. Nodes move among healthy, compromised, detected, isolated, restoring, and restored states. Propagation can occur only along a directed edge and is modified by vulnerability, patch status, privilege, segmentation, identity compromise, backup access, and the selected portfolio.

Three neutral network scales were used in a cross-setting replication. Their node counts were 40 to 60, 200 to 300, and 750 to 1,000. Three synthetic postures were also tested. These labels do not correspond to real hospital size, geography, or maturity. Entry types were a workstation, internet-facing service, privileged system, medical device, or vendor connection.

2.2 Service outcomes

The primary outcome was weighted service-hours lost across EHR, laboratory, pharmacy, imaging, scheduling, identity, and backup/recovery. At each step, a service

was unavailable if its core was nonfunctional, fewer than 60% of its support nodes remained functional, or a required identity dependency was unavailable. Functional means neither actively compromised nor isolated, so defensive isolation can itself remove service capacity.

Reference weights were 1.0, 0.8, 0.9, 0.7, 0.4, 0.9, and 0.6, respectively. The 60% rule and weights are value judgments, not clinical standards. We therefore retained unweighted, service-specific downtime in every row and varied the functional threshold from 45% to 75% in joint stress tests. A secondary sustained-outage indicator was set when any of four clinical services was unavailable for more than two modeled hours. We avoid the earlier term “catastrophic risk” because the indicator is neither clinically validated nor a real-world probability.

2.3 Evidence and parameter provenance

Every influential input was classified as empirical observation, evidence-informed window, outcome definition, or stress assumption. Public studies informed the 72-hour acute window and the need for 7- and 21-day recovery checks. Mechanistic values such as per-edge spread, patch multiplier, detection delay, isolation success, identity escalation, restoration throughput, false-positive isolation, legacy fraction, and residual backup access were not identifiable from those studies. They remained synthetic stress variables.

The final joint design varied 13 assumptions with Latin-hypercube sampling. The range for 15-minute-equivalent spread was 0.05 to 0.50. Patch effectiveness ran from 0.30 to 0.95. Within-zone degree ranged from 2 to 5, while cross-zone degree ranged from 1 to 4. We also varied identity escalation, restoration, backup access, the service threshold, false isolation, legacy assets, containment, and detection time. The complete ranges are recorded in the parameter register. Clock-time hazards and durations were converted for the five-minute replication. These are stress ranges, not estimates for the hospital sector.

2.4 Observed datasets and linkage rules

We kept observed records outside the synthetic scenario banks. The first source was the public THREAT hospital/market event file deposited with the replication materials for Neprash et al. [5] [18]. It contains Medicare identifiers, attack dates, event identifiers, emergency-diversion and cancellation/delay flags, nearest-neighbor indicators, and up to seven origin event identifiers. The public preview contained 1,999 rows. We trimmed display whitespace, retained all 14 columns, and collapsed duplicate attacked hospital-event rows by Medicare identifier, THREAT identifier, and attack date, taking the maximum of duplicate operational flags. The resulting analysis set contained 149 attacked hospital-event records, 148 distinct hospitals, and 74 events from 2016 through 2021. This is organization-level public information, not patient-level data.

The second source was CISA’s Known Exploited

Vulnerabilities catalog, frozen at version 2026.09.01 [19]. We retained the catalog metadata and all 1,687 records, then selected the 352 entries whose `knownRansomwareCampaignUse` field was `Known`. Vendor, product, CVE, date, and required-action fields provide auditable provenance for future entry scenarios. Catalog membership does not measure hospital exposure, attack frequency, exploit success, patch effectiveness, or the probability that ransomware caused a healthcare incident.

Before looking for numerical similarity, we classified each observed-to-model mapping as a direct structural check, a low-comparability operational proxy, a construct bridge, a duration-pattern check, or scenario provenance only. We did not fit any transition, detection, containment, or recovery parameter to these records, and we did not generate new simulations. The purpose was to identify where an existing output has a defensible observational interpretation and where the model lacks the construct entirely.

2.5 Portfolios and pairing

The reference used a flat graph, the posture's baseline patch and response values, and connected backups. The layered portfolio combined basic segmentation, a one-tier detection improvement, rapid same-step isolation attempts, and protected backups. "Protected backup" is an abstraction that compresses credential separation, management-plane access, immutability, testing, reconnection, and reinfection into a small set of paths. Residual access was explicitly stressed.

All comparisons used common random numbers, but not by simply restarting a stateful random-number generator with the same seed. That shortcut can lose event alignment when an intervention changes the code path [20]. Instead, a stable key made from the scenario, event type, and time step generated a full node- or edge-level random vector. A portfolio could change which entries were used, but it could not shift the values assigned to events shared with the other portfolio. The generated base topology, entry type and node, and patch draws were also held fixed within a pair. Thus a portfolio did not receive an easier network simply by chance.

2.6 Study sequence and sample size

Earlier repository results were treated as development evidence and were not pooled with the new study. Before generating the 15-minute validation, we froze a local protocol, configuration, seeds, and file hashes. Planning data gave a paired-difference standard deviation of 73.90 service-hours and a reference mean of 169.97. Requiring Monte Carlo standard error no greater than 2% of that reference mean gave 473 pairs. We rounded to 500.

The 15-minute phase contained 500 primary pairs, 100 scenarios in each of nine scale-by-posture combinations, and 128 joint stress settings with 20 scenarios each. Pre-specified 5-, 15-, and 30-minute diagnostics initially revealed a scenario-pairing and clock-time conversion er-

ror. Original files were retained, the deviation was documented, and only those diagnostics were corrected. The corrected diagnostic then showed material time-step sensitivity. An explicitly exploratory 1-, 2-, and 5-minute pilot was run. Finally, a post-diagnostic five-minute protocol was frozen before 500 new primary pairs and 128 new stress settings with 10 scenarios each. It used fresh seeds and scenario identifiers.

The new study contains 19,340 executions. This total includes 11,880 runs in the 15-minute study, 1,800 retained runs from the flawed diagnostic, 1,800 corrected diagnostic runs, 300 exploratory pilot runs, and 3,560 runs in the frozen five-minute replication. Superseded rows remain in the archive and were never pooled with corrected results. Here, a run means one portfolio under one synthetic scenario. It does not mean one hospital or one real incident.

2.7 Analysis and verification

The primary statistic was the within-scenario difference (layered minus reference) in weighted service-hours. We report means, medians, the distribution of paired changes, and 95% percentile bootstrap intervals that re-sample complete scenarios. Joint-stress results were first averaged within parameter setting. Robustness required a negative primary interval, lower mean loss in at least 95% of settings, and a positive fifth percentile of setting-specific relative reduction. This success rule concerns model behavior only.

Fifty-nine unit and invariant tests and twelve behavioral validation cases covered transitions, pairing, seed reproduction, segmentation, patching, isolation, backup behavior, dependencies, budgets, export, observed-data deduplication, and source-count consistency. An independent script recomputed all three derived ratios from raw numerators, denominators, and final-service flags. The study follows the transparency intent of the STRESS reporting guideline [21]. We retained the protocols, parameter registers, deviations, raw outputs, and analysis manifests.

3 Results

3.1 Five-minute paired replication

Across 500 fresh pairs, reference scenarios averaged 193.4 weighted service-hours lost. Layered scenarios averaged 33.9. The paired mean change was -159.5 hours, with a 95% bootstrap interval from -166.1 to -152.7 . Its Monte Carlo standard error was 3.42, below the protocol target of 3.87. Loss was lower in 451 scenarios and tied in 48. One scenario was worse by 11.35 hours. The mean-of-groups reduction was 82.5% *within the synthetic model*. It is not a real control-effectiveness estimate.

The reference sustained-outage indicator was positive in 90.2% of scenarios and the layered indicator in 41.2%, a paired difference of -49.0 percentage points. These are conditional model frequencies under a deliberately severe scenario bank, not incident probabilities.




Figure 1: Primary five-minute results. The first two panels compare the group means. The third shows the direction of all 500 paired comparisons. The mean paired change was -159.5 weighted service-hours, and its 95% bootstrap interval ran from -166.1 to -152.7 . All values are conditional on the synthetic model and its five-minute step.

Service-specific results did not depend on weights for their direction. Mean EHR downtime changed from 32.9 to 7.4 hours, laboratory from 36.8 to 7.3, pharmacy from 34.7 to 6.9, imaging from 40.7 to 7.6, scheduling from 49.3 to 8.0, identity from 32.1 to 6.2, and backup/recovery from 38.0 to 0.6 *within the synthetic model*. These values describe binary modeled availability, not hours of complete clinical shutdown.




Figure 2: Mean modeled downtime for each service in the 500 fresh pairs. Every service had a lower mean under the layered portfolio. These are binary service states in a synthetic model, not measured hours of hospital shutdown.

3.2 Joint uncertainty

The layered portfolio had lower mean loss in all 128 five-minute joint stress settings. The Wilson 95% interval for that fraction ran from 97.1% to 100%. The median setting-specific reduction was 28.7%, the fifth percentile was 3.1%, and the 95th percentile was 96.4%. Figure 4




Figure 3: Distribution of paired five-minute effects. Negative values favor the layered portfolio. Most scenarios improved substantially, 48 tied at zero, and one was worse. The complete distribution keeps the mean from hiding that exception.

shows the full distribution. The success rule was met, but the low fifth percentile means the advantage was small in some plausible stress combinations. In the earlier 15-minute stress bank, the median and fifth percentile were 66.5% and 39.0%, respectively. The direction replicated. The magnitude did not.

3.3 Cross-setting and numerical checks

In the independent 15-minute cross-setting bank, the layered mean was below the reference mean in all nine scale-by-posture cells. Reductions of the cell means ranged from 55.2% to 94.2%. Scenario-level improvement fractions were lower in stronger-posture cells because many reference scenarios had little or no loss, creating ties. These labels remain synthetic. The result does not show that a particular real hospital category would benefit by those percentages.

The corrected coarse-step diagnostic produced lay-




Figure 4: Five-minute joint stress results. Every setting favored the layered portfolio, but reductions ranged from nearly zero to very large. The lower tail, not the best case, limits the strength of the conclusion.

ered/reference reductions of 81.7%, 88.4%, and 94.0% at 5, 15, and 30 minutes. The exploratory fine-step pilot produced 82.3%, 82.6%, and 85.6% at 1, 2, and 5 minutes. Paired mean differences in that pilot ranged from -155.8 to -175.0 hours with Monte Carlo standard errors near 10 hours. Thus, fine resolutions gave a broadly similar comparative answer, but neither absolute values nor exact percentages were numerically invariant. The 30-minute step was especially optimistic for the layered portfolio because fewer propagation generations can occur before a fixed clock-time response.

At the 15-minute reference settings, 95% of flat-reference scenarios recovered within 72 hours and 100% within seven days. This is shorter than the multi-week disruptions documented in public incident reports [1]. The model therefore fails that empirical duration pattern under its reference recovery values. We do not interpret modeled recovery time as a prediction.

3.4 Observed-data benchmark

The 149 deduplicated attacked hospital-event records included 46 with an emergency-diversion flag (30.9%), 73 with a cancellation/delay flag (49.0%), and 85 with either flag (57.0%). Nineteen of 74 events (25.7%) affected more than one Medicare hospital in the file; the largest event covered 22. These descriptive results apply to the observed, publicly documented and Medicare-linkable event set. They are not sector probabilities.

Table 1 records the external checks. The diversion and cancellation fields support the decision to measure services rather than only compromised assets, but they cannot calibrate the sustained-outage indicator. The event-breadth comparison is stronger: one facility per run is structurally incapable of reproducing the multi-hospital events in the data. The published hospital-volume result supplies another direct warning about duration: volume fell 17% to 24% in the first attack week and recovered over roughly three weeks [5], whereas every reference case in our 7-day diagnostic had technically recovered. These are different recovery constructs, and the modeled process is too fast even before that distinction is considered.

3.5 Derived-metric reconciliation

Anne Lee’s review noted that three secondary ratios had been excluded without a satisfactory explanation. Independent reconstruction showed that the archived `pct_compromised` field equals ever-compromised nodes divided by total nodes, and `pct_clinical_capacity_lost` equals time-averaged weighted unavailability of four clinical services. Both agreed across all 15 archived CSV files to within 5.0×10^{-11} .

The third name, `pct_services_restored`, was misleading. It measured the fraction of all seven services available at the final step, including services never disrupted. Legacy rows did not store the seven final flags, so the field was not independently row-auditable. Exact reruns of sampled archived rows still matched. New rows retain the historical alias for compatibility but add explicit final flags, `final_service_availability_fraction`, and `disrupted_service_recovery_fraction`. Independent recomputation of all new validation files agreed to within 5.0×10^{-11} . No headline analysis uses the ambiguous legacy names.

4 Discussion

4.1 What the comparison showed

Our first result looked almost too clean. At a 15-minute step, the layered portfolio’s median reduction across stress settings was 66.5%. Repeating the study at five minutes left the direction unchanged. All 128 settings still favored the layered portfolio. The median fell to 28.7%, and the lower tail was close to zero. That diagnostic changed how we interpret the paper. The 82.5% reduction in the primary scenario bank is a description of those runs, not a stable estimate of control effectiveness.

The direction is less surprising than the amount. We encoded segmentation to remove paths, faster response to shorten the time available for spread, and protected backups to preserve a recovery route. The useful observation is that their combined advantage remained after 13 assumptions were varied together and after the simulation step was shortened. This agrees with layered-defense guidance [6–8]. It does not show that our numerical multipliers match controls deployed in a hospital.

4.2 How the workflow could be used

We see a possible use in tabletop planning, but only after external review and local calibration. A hospital could describe service dependencies without recording sensitive architecture, replace our synthetic ranges with authorized summaries, and define degraded operation with clinical staff. The output should be a short robustness report: which portfolios keep their ranking, which assumptions reverse it, and which missing measurements matter most. It should not be treated as a substitute for a risk assessment, architecture review, incident-response plan, regulatory obligation, or procurement process.




Figure 5: The estimated reduction changed with the simulation step. Values within each panel share a scenario bank, but the exploratory pilot and the corrected diagnostic use different scenario banks. The direction stayed the same while the size changed.

Table 1: Observed-to-simulation bridge. Observational percentages describe the public source records, not hospital-sector incidence. “Fail” denotes a missing or contradicted structure, not a failed statistical goodness-of-fit test.

Observed record	Existing model output	Interpretation
Emergency diversion: 46/149 (30.9%) Cancellation/delay: 73/149 (49.0%)	Sustained outage: 90.2% reference, 41.2% layered Seven binary service-downtime outputs	Low comparability; demand, staffing, policy, and regional capacity are absent, so no calibration was performed. Construct bridge only; modeled hours cannot be converted to canceled cases.
Multi-hospital events: 19/74 (25.7%)	Exactly one facility per run	Direct structural fail; multi-facility spread and regional spillover are impossible.
Volume down 17–24% in week 1; recovery about 3 weeks	100% reference recovery within 7 days	Duration-pattern fail; technical and operational recovery also differ in definition.
CISA: 352/1,687 KEVs tagged for ransomware use	Generic vulnerability and patch scalars	Scenario provenance only; the catalog share is neither an attack probability nor patch efficacy.

4.3 Where the model breaks down

The real-data layer measures event timing and breadth, two operational flags, and ransomware-linked vulnerability identities. It does not measure hospital topology, permitted paths, attack transitions, detection and containment, restoration capacity, service dependencies, deployed control state, or control costs in the same incidents. It therefore reduces unsupported statements about what was observed but does not identify the model’s mechanistic parameters. The confidence intervals cover Monte Carlo variation inside the model and nothing more. They do not quantify measurement error, selection, structural error, or transferability to a hospital.

The THREAT file contains publicly documented, Medicare-linkable records rather than a random hospital sample. A zero flag can mean that diversion or cancellation was not reported, not necessarily that it did not occur. The committed CSV is normalized from the public preview because original-file download requires an ICPSR sign-in. Before submission, the depositor’s original CSV and codebook should replace it, with row and value reconciliation retained as a documented transformation. The CISA catalog is also curated: it contains known exploited vulnerabilities, not the denominator of vulnerabilities on a hospital’s assets. Its 20.9% ransomware-

tagged share must not be interpreted as a base rate.

The time-step diagnostic is the clearest technical weakness. Finer steps allow more rounds of propagation before the same clock-time response, and they change the order of events. Five-minute runs reduced the discrepancy but did not make the absolute results invariant. We would replace this mechanism with an event-driven continuous-time model before interpreting exact hours or percentages.

Service availability is binary and depends on one support threshold. Real departments operate in degraded modes using paper, phones, manual verification, reduced menus, and workarounds whose effectiveness changes with time and staff load [9, 10, 22]. Because our model does not contain patient flow, the phrase “service-hours lost” cannot be translated into patients delayed, diverted, or harmed. Regional spillover observed by Dameff et al. [3] and Abouk and Powell [17] is also absent.

Synthetic postures bundle patching, detection, isolation, legacy assets, segmentation, and backups. They are scenario devices, not validated maturity classes. “Protected backups” simplify credentials, management planes, immutability, testing, reconnection, and reinfection. Normalized budget points are not dollars. The port-

folio search cannot discover a control or failure mode that was not encoded.

Public reports create another problem: visible incidents are selected, and “recovery” may refer to technical restoration, clinical workflow, billing, or normal operations. Even with that caution, our reference model recovered far sooner than the multi-week pattern reported in the literature. We treat that mismatch as evidence that its recovery output should not be used predictively.

4.4 Next validation step

This paper stops before outside contact. The next step is a structured review by hospital IT or incident-response staff and by clinical-operations staff, with privacy and research-ethics preparation completed first. We would ask about state transitions, degraded service, dependencies, timing, backup failure modes, and useful outputs. We would not request network diagrams, credentials, vulnerability details, patient data, or active-incident information. A review of this kind would establish face validity at most. Predictive validation would still require suitable withheld incident or exercise data.

5 Conclusion

In our fresh paired runs, basic segmentation, faster detection and isolation, and protected backups reduced modeled service disruption relative to a flat reference. The ranking survived every joint stress setting at both 15- and 5-minute resolution. The estimated size did not. At five minutes, some stress settings showed only a small advantage, and absolute outcomes still changed with numerical resolution.

The observed-data bridge narrows what that result can mean. It shows that the current model cannot represent multi-hospital events, cannot turn digital downtime into diversion or canceled care, and recovers too quickly relative to the published hospital-volume pattern. CISA records give future entry scenarios real vulnerability identities but do not supply probabilities or effect sizes.

We are therefore confident about one result inside this model: the layered portfolio was directionally robust over the uncertainty we tested. We are not confident that the hours or percentages transfer to a hospital. The metric repair, retained diagnostic error, and failed recovery-duration check are part of the result because they show where the analysis stops being reliable. Operational review and a continuous-time model with degraded-service states are needed before using the framework for a real decision.

Research integrity and declarations

Ethics. The study used synthetic networks, simulated trajectories, public organization-level hospital-event records, a public vulnerability catalog, and published aggregate findings. It involved no participants, patient-level records, private organizational data, mal-

ware, or intervention. External interviews have not been conducted. Consent, privacy, and institutional requirements must be checked before collecting information from outside reviewers.

Data and code. The revised code, protocols, configurations, manifests, raw simulation outputs, public observed-data snapshots, processed summaries, source hashes, tests, figures, and manuscript source are at <https://github.com/Dodgeboi/NazeResearchNSRI>. The development history is retained in that repository. Any submitted version should cite a fixed release or archived commit.

Artificial-intelligence use. We used Claude, ChatGPT, and Codex throughout the project. We asked them for coding help, debugging ideas, literature search terms, analysis checks, organization, and editing suggestions. We did not use their responses as evidence. We checked numerical claims against the code and CSV files, and we checked literature claims against the cited sources. This process helped us find the time-step problem. We kept that problem in the study record instead of hiding it. Each author is responsible for the final analysis and wording.

Author contributions. Ashish Agrawal, Mukil Dharanidharan, and Naman Upadhyay worked together on the project and reviewed this manuscript. The authors will agree on detailed CRediT roles before a formal submission.

Funding. The authors report no specific funding for this work.

Conflicts of interest. The authors report no known competing interests. All authors should reconfirm this statement before submission.

References

- [1] Hannah T. Neprash, Claire C. McGlave, David A. Cross, et al. Trends in ransomware attacks on US hospitals, clinics, and other health care delivery organizations, 2016–2021. *JAMA Health Forum*, 3(12):e224873, 2022. doi: 10.1001/jamahealthforum.2022.4873.
- [2] Saira Ghafur, Soren Kristensen, Kate Honeyford, Guy Martin, Ara Darzi, and Paul Aylin. A retrospective impact analysis of the WannaCry cyberattack on the NHS. *npj Digital Medicine*, 2:98, 2019. doi: 10.1038/s41746-019-0161-6.
- [3] Christian Dameff, Jeffrey Tully, Theodore C. Chan, et al. Ransomware attack associated with disruptions at adjacent emergency departments in the US. *JAMA Network Open*, 6(5):e2312270, 2023. doi: 10.1001/jamanetworkopen.2023.12270.
- [4] Thaidan T. Pham, Theoren M. Loo, Atul Malhotra, Christopher A. Longhurst, Diana Hylton, Christian Dameff, Jeffrey Tully, Gabriel Wardi, Rebecca E. Sell, and Alex K. Pearce. Ransomware cyberattack associated with cardiac arrest incidence and outcomes at untargeted, adjacent hospitals. *Critical Care Explorations*, 6(4):e1079, 2024. doi: 10.1097/CCE.0000000000001079.
- [5] Hannah T. Neprash, Claire C. McGlave, and Sayeh Nikpay. Hacked to pieces? the effects of ransomware attacks on hospitals and patients. *American Economic Journal*:

- Economic Policy*, 18(1):256–281, 2026. doi: 10.1257/pol.20240594.
- [6] Cybersecurity and Infrastructure Security Agency, Federal Bureau of Investigation, National Security Agency, and Multi-State Information Sharing and Analysis Center. Stopransomware guide. Technical report, CISA, 2023. URL <https://www.cisa.gov/resources-tools/resources/stopransomware-guide>.
- [7] U.S. Department of Health and Human Services. Healthcare and public health sector cybersecurity performance goals, 2026. URL <https://hhs cyber.hhs.gov/cybersecurity-performance-goals.html>. Accessed July 18, 2026.
- [8] William Fisher, Murugiah Souppaya, William Barker, and Karen Kent. Ransomware risk management: A cybersecurity framework 2.0 community profile. Technical Report NIST IR 8374 Rev. 1, National Institute of Standards and Technology, 2026.
- [9] Ethan Larsen, Allan Fong, Christian Wernz, and Raj M. Ratwani. Implications of electronic health record downtime: An analysis of patient safety event reports. *Journal of the American Medical Informatics Association*, 25(2):187–191, 2018. doi: 10.1093/jamia/ocx057.
- [10] Ethan Larsen, Daniel Hoffman, Christian Rivera, Brian M. Kleiner, Christian Wernz, and Raj M. Ratwani. Continuing patient care during electronic health record downtime. *Applied Clinical Informatics*, 10(3):495–504, 2019. doi: 10.1055/s-0039-1692678.
- [11] Christian Wernz. Evidence-based contingency planning for electronic health record downtime: Final report. Technical Report R21 HS024350, Agency for Healthcare Research and Quality, 2018. URL <https://digital.ahrq.gov/>.
- [12] Hadi Ghayoomi, Kathryn Laskey, Elise Miller-Hooks, Charles Hooks, and Mersedeh Tariverdi. Assessing resilience of hospitals to cyberattack. *Digital Health*, 7, 2021. doi: 10.1177/20552076211059366.
- [13] Markus Willing, Simon Ebberts, Christian Dresen, Marc Czolbe, Christoph Saatjohann, and Sebastian Schinzel. Simulating the overload of medical processes due to system failures during a cyberattack. *BMC Medical Informatics and Decision Making*, 2025. doi: 10.1186/s12911-025-02988-8.
- [14] Stefano Chiaradonna, Petar Jevtić, and Nicolas Lanchier. Framework for cyber risk loss distribution of hospital infrastructure: Bond percolation on mixed random graphs approach. *Risk Analysis*, 2023. doi: 10.1111/risa.14127.
- [15] David Carramiñana, Ana M. Bernardos, Juan A. Besada, and José R. Casar. Enhancing healthcare infrastructure resilience through agent-based simulation methods. *Computer Communications*, 234:108070, 2025. doi: 10.1016/j.comcom.2025.108070.
- [16] Pouria Arsalani, Iman Dayarian, Nickolas Freeman, Youngsoo Kim, David Peng, Ming Zhao, Jamie Gray, and William Crawford. Impact of ransomware on emergency medical services: Empirical estimation and risk assessment framework. *Omega*, page 103578, 2026. doi: 10.1016/j.omega.2026.103578.
- [17] Rahi Abouk and David Powell. Ransomware attacks, emergency department visits and inpatient admissions in targeted and nearby hospitals. *JAMA*, 331(24):2129–2131, 2024. doi: 10.1001/jama.2024.7752.
- [18] Hannah Neprash, Claire McGlave, and Sayeh Nikpay. Data and code for: Hacked to pieces? the effect of ransomware attacks on hospitals and patients, 2025. CC BY 4.0.
- [19] Cybersecurity and Infrastructure Security Agency. Known exploited vulnerabilities catalog, 2026. URL <https://www.cisa.gov/known-exploited-vulnerabilities-catalog>. Frozen JSON snapshot released September 1, 2026.
- [20] Vince Buffalo, Carl A. B. Pearson, and Daniel Klein. Realizing common random numbers: Event-keyed hashing for causally valid stochastic models, 2026.
- [21] Thomas Monks, Christine S. M. Currie, Bhakti S. Onggo, Stewart Robinson, Martin Kunc, and Simon J. E. Taylor. Strengthening the reporting of empirical simulation studies: Introducing the STRESS guidelines. *Journal of Simulation*, 13(1):55–67, 2019. doi: 10.1080/17477778.2018.1442155.
- [22] Joseph M. Walsh, Elizabeth M. Borycki, and Andre W. Kushniruk. Strategies in electronic medical record downtime planning: A scoping study. *Studies in Health Technology and Informatics*, 257:449–454, 2019.